

Model Provider	Model	Release Date	Cut Off Date	Cost Of Tokens	Context Window	Model Type	When to use this model
Open AI	GPT-5.5 (Flagship)	April 23, 2026	December 1, 2025	\$5 per 1M Input and \$30 per 1M Output tokens	1M	Public	Designed specifically for agentic tasks, complex reasoning, and high-stakes coding rather than general chat or simple summarization. It is optimized for multi-step, tool-using workflows where the model needs to plan, act, and self-correct with minimal human intervention.
Google (Gemini)	Gemini 3.1 Pro	February, 2026	January, 2025	\$2 per 1M Input and \$12 per 1M Output tokens	1M	Public	For complex, high-reasoning tasks requiring advanced logic, such as software engineering, analyzing large datasets (1M+ token context), or multi-modal understanding across text, code, images, and video. It is best suited for agentic workflows, deep research, and creating complex, interactive applications.
Anthropic	Claude Mythos	May, 2026	December, 2025	\$25 per 1M Input and \$125 per 1M Output tokens	1M	Private	Claude Mythos is a highly advanced, gated AI model from Anthropic designed for elite cybersecurity, software engineering, and defensive "red teaming" tasks. Due to its immense power to both find and fix zero-day vulnerabilities, it is used exclusively by a limited group of security researchers
Meta (AI)	Muse Spark	April, 2026	No Cutoff officially	\$0.0 per 1M Input and \$0.0 per 1M Output tokens	262144	Public	Used for complex, multimodal reasoning tasks, such as analyzing images/charts, advanced scientific/health inquiries, and generating interactive content (minigames, visual troubleshooting). It excels in its "Contemplating Mode," which uses multiple agents to provide deeper, more reliable answers for research and coding, acting as a strong, costeffective alternative to top-tier models.
NVIDIA	Nemotron 3 family	December 15, 2025	November, 2025	\$0.04 to \$0.10 per 1M input and \$0.16 to \$0.50 per 1M tokens	1M	Public	Used for building highly efficient, agentic AI systems that require fast, multi-step reasoning, tool integration, and long-context understanding (up to 1M tokens). The models use a hybrid Mamba-Transformer MoE architecture to provide superior throughput over standard transformers.

Microsoft Azure AI	MAI-Transcribe-1, MAI-Voice-1, and MAI-Image-2	April, 2026	January, 2026	\$5 per 1M Input and \$33 per 1M Output tokens	NA	Public	MAI-Transcribe-1 (Speech-to-Text): Designed for long-form audio, it accepts MP3, WAV, and FLAC files, with specific capabilities to process lengthy recordings for meeting transcription and subtitles. While not explicitly defining a "token" context window size like an LLM, it is optimized to handle large, noisy audio files. MAI-Voice-1 (Text-to-Speech): This model is designed for high-fidelity audio generation, specifically noted for preserving speaker identity across long-form content. It can generate 60 seconds of audio in under one second. MAI-Image-2 (Text-to-Image): This model generates images from text prompts and supports detailed, in-image text rendering. It is geared toward high-fidelity visual generation.
Amazon Web Services (AWS)	Amazon Nova family	May, 2026	December, 2024	\$0.035 per 1M input and \$0.14 per 1M output	128000	Public	Designed for enterprise-scale AI applications that require a balance of high performance, multimodal capabilities, and cost efficiency. You should use the Amazon Nova suite when building production-ready generative AI applications—such as chatbots, content generators, or intelligent agents—that need to process text, images, or video within the AWS ecosystem.
ElevenLabs	Eleven v3	May, 2026	March, 2026	1 credit per character	5,000 characters per request	Public	Used for high-expressivity, long-form content, such as audiobooks, dramatic narrations, or complex, multi-speaker dialogue requiring emotional nuance.
DeepSeek	DeepSeek-V4 Pro and DeepSeek-V4 Flash	April, 2026	July, 2026	\$1.74 per 1M input and \$3.48 per 1M output tokens	1M	Public	For advanced reasoning and complex workflows
Salesforce	Agentforce 360 Platform	October, 2025	December, 2025	\$10 per 1M input and \$40 per 1M output tokens	Agentforce manages context windows by using "compaction" to summarize conversation history, ensuring agents stay on track without hitting token limits.	Public	Use Agentforce when you need to automate complex, multi-step business processes within Salesforce, moving beyond simple chatbots to autonomous agents that reason, plan, and take action. It is best applied for 24/7 customer service, high-volume lead qualification, personalized marketing, and providing employees with intelligent, data-driven assistance.